Efficient Deep Learning Systems
Course introduction

Max Ryabinin

2025

What's this about?

- DL as a field is getting mature:
  - Neural networks are becoming more and more widespread in practice
  - Scaling trends everywhere (model size, dataset size, coauthor list size)
  - .ipynb-based development is no longer sustainable :)
- Each model is much more than just architecture, loss and even data
- Engineering knowledge becomes handy even for SOTA research
- For practical applications, performance and maintainability are key factors

Bird's eye view of DL

Training	→	Inference
How to achieve the best quality?		~~Is my model useful?~~ ~~Is my model good enough?~~
Do I utilize my resources to the fullest?		Is performance good enough for my use case?
How to navigate 100s of experiments?		How do I ensure the model is maintainable?
How to avoid bugs in my pipeline?		How to avoid bugs in my pipeline? pt.2

Goal of the course

- Most DL courses do not cover practical details and overall systems:
  - Small code changes can make your training/inference much faster
  - Deployment of trained networks, both on their own and as a part of a larger system
  - Streamlined maintenance by treating ML models like any other code (testing, versioning, etc.)
- Knowledge about this is scattered around the Internet and unstructured
- We want to give you these useful bits of practical knowledge!
- …no bleeding-edge methods or last-week papers (with some exceptions)

Plan

1. *(You are here)* Intro, basics of GPU architecture & benchmarking
2. Experiment tracking & versioning, testing & debugging
3. Profiling DL pipelines, tricks for efficient training
4. Data-parallel training, All-Reduce, torch.distributed intro
5. Sharded data parallel training, optimization tips
6. Memory-efficient training, model parallelism
7. Basics of web service deployment
8. Deploying neural networks: software side
9. Optimizing models for inference

Systems & better training 1

Basically, MLOps

Systems & better training 2

Distributed training

Deployment in production

Logistics

- Lectures&seminars: every Wednesday, 18:00 – 21:00, via Zoom
- Course repo: github.com/mryab/efficient-dl-systems
- Anytask/LMS for handing in assignments
- Channel with announcements: see HSE FCS wiki/course page in LMS
- Resources: Yandex Cloud VM + DataSphere (HSE), YSDA GPUs + DataSphere (YSDA)

Grading

- 3 assignments:
    1. Fast and reliable pipelines (2 parts)
    2. Distributed training (3 parts)
    3. Deployment (3 parts)

- Each assignment consists of sub-assignments given each week (except this one)
- Final grade: $G_{total} = 0.2G_1 + 0.4G_2 + 0.4G_3$

GPU architecture: a brief overview

- As the name suggests, originally used for graphics
- Highly parallel execution model: objects can be rendered simultaneously
- Since ~2007, simple GPGPU API started to appear (CUDA, OpenCL, Metal)
- GPU-trained AlexNet/DanNet sparked the DL revolution in early 2010s

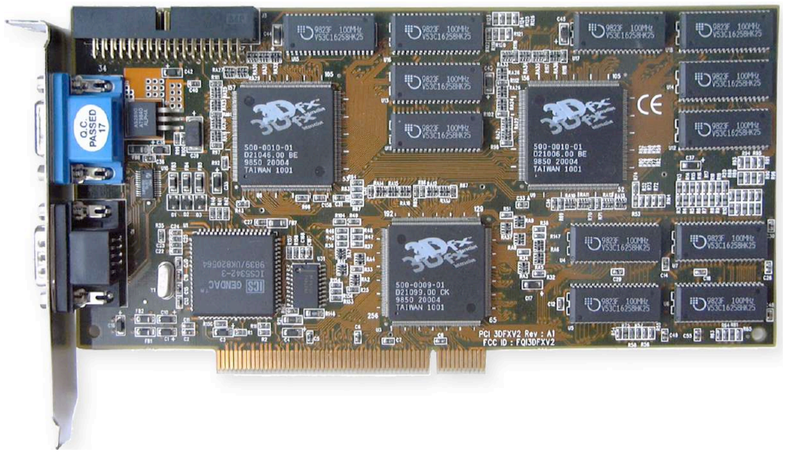

3dfx Voodoo2: 12MB RAM

NVIDIA H200: 141GB RAM

GPU architecture: a brief overview

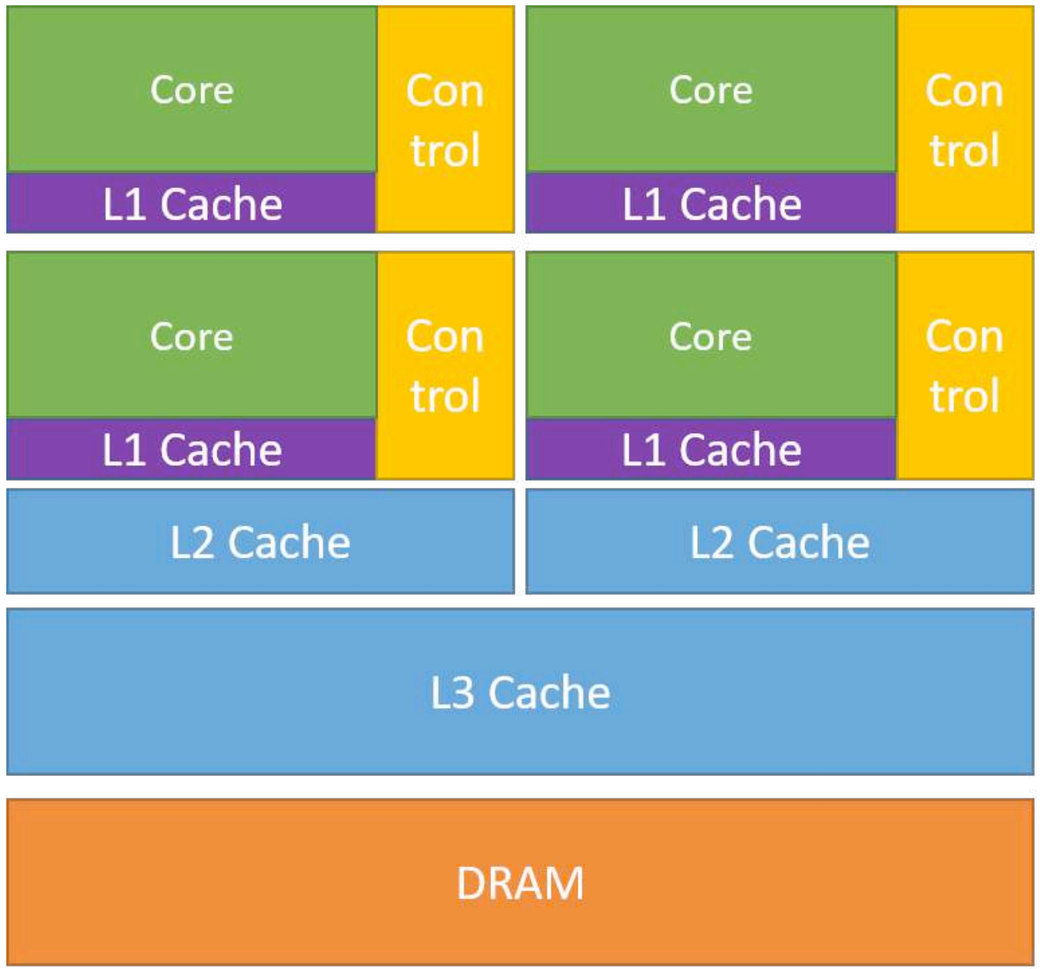

CPU

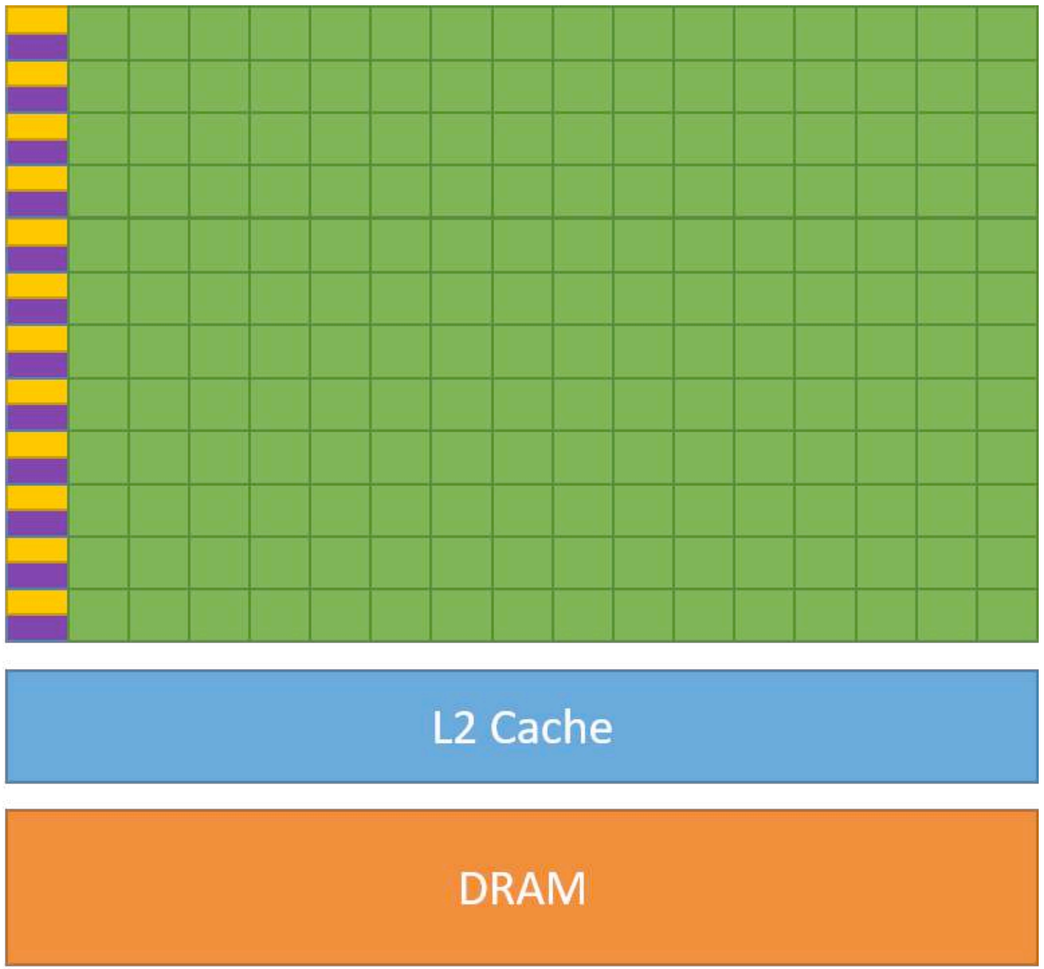

GPU

docs.nvidia.com/cuda/cuda-c-programming-guide

CUDA computation model

- In CUDA, we launch kernels from the host that are executed in parallel on the device
- Kernels are executed by threads grouped in thread blocks of limited size
- A GPU is composed of multithreaded Streaming Multiprocessors (SMs) that are assigned different thread blocks
- Multiple thread blocks are arranged in grids (can be 1D, 2D, or 3D)

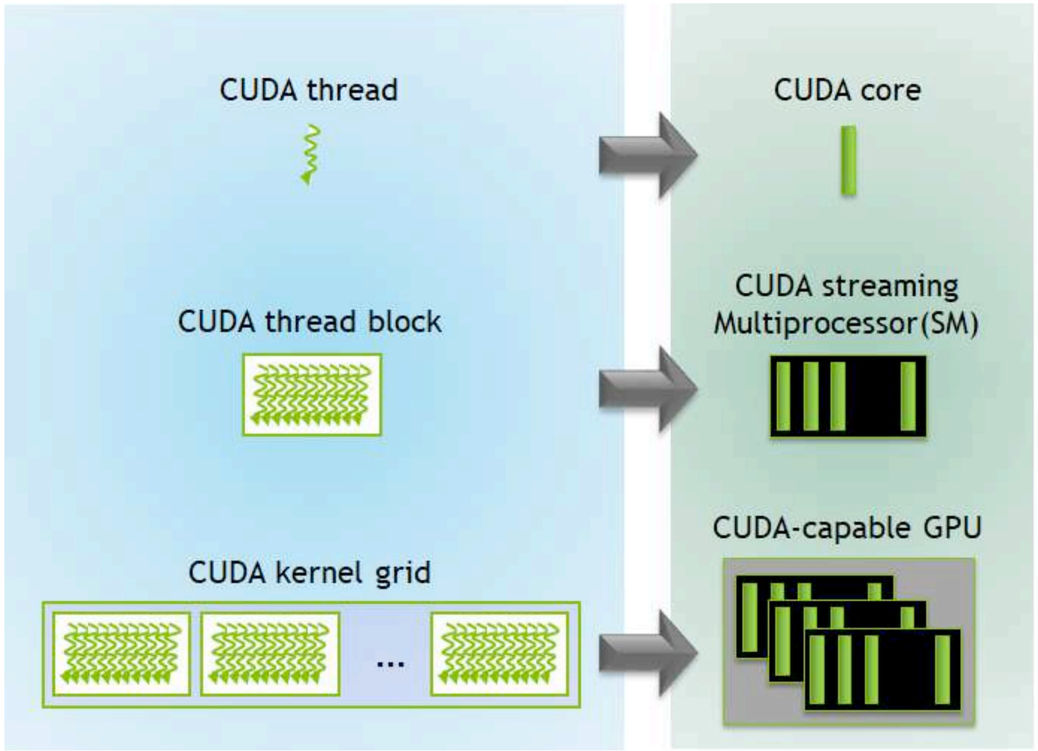

developer.nvidia.com/blog/cuda-refresher-cuda-programming-model

GPU computations: hardware side

- SIMT (Single Instruction, Multiple Thread)
- On a physical level, threads are executed in groups of 32 called warps
- A warp executes one instruction at a time:
  in case of branching, all paths need to be taken
- This does not affect correctness but has major performance implications
- Warp-level primitives can be leveraged for parallel computation

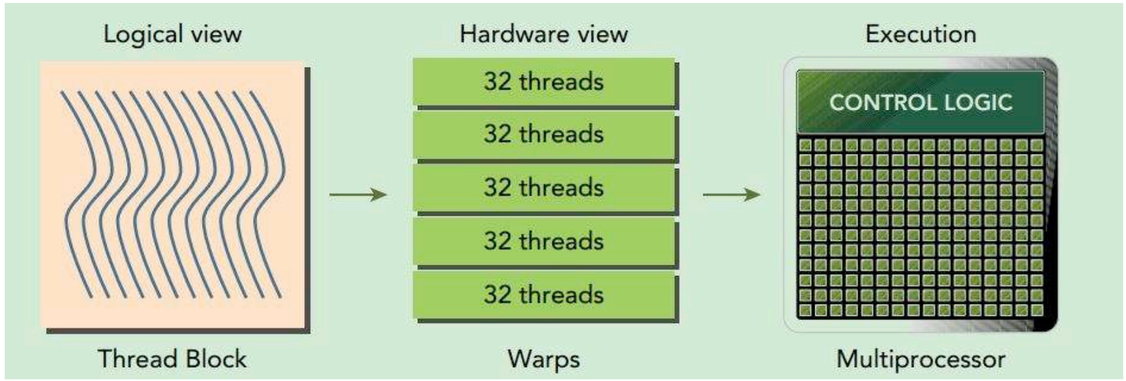

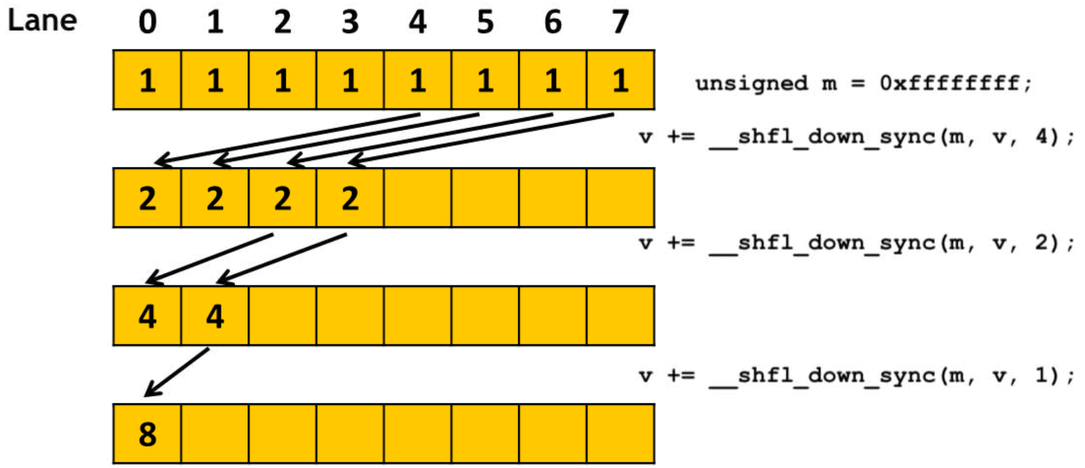

developer.nvidia.com/blog/using-cuda-warp-level-primitives

Why does all of this matter?

- The most popular operation in DL is matrix multiplication
- Executing this in parallel can have two potential effects when dividing the work
- **Tile Quantization**: matrix size is not divisible by the thread block tile size

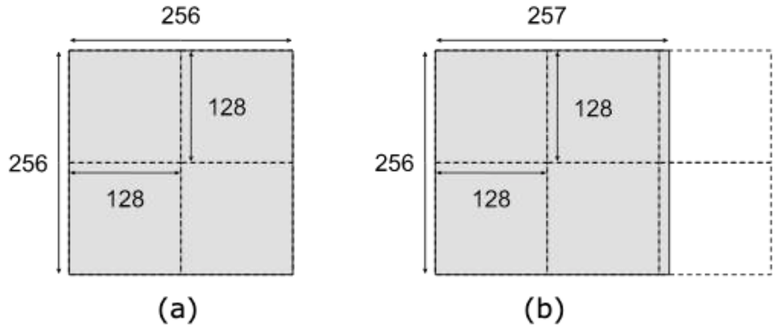

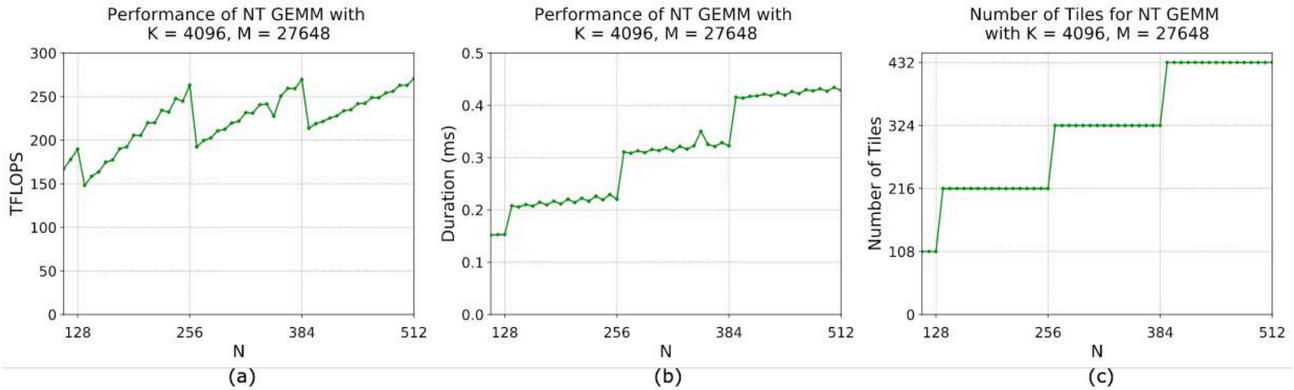

docs.nvidia.com/deeplearning/performance/dl-performance-matrix-multiplication/index.html

Why does all of this matter?

- The most popular operation in DL is matrix multiplication
- Executing this in parallel can have two potential effects when dividing the work
- **Tile Quantization**: matrix size is not divisible by the thread block tile size
- **Wave Quantization**: total number of tiles is quantized to the number of SMs
- Both effects can be quite noticeable for small or irregular shapes!

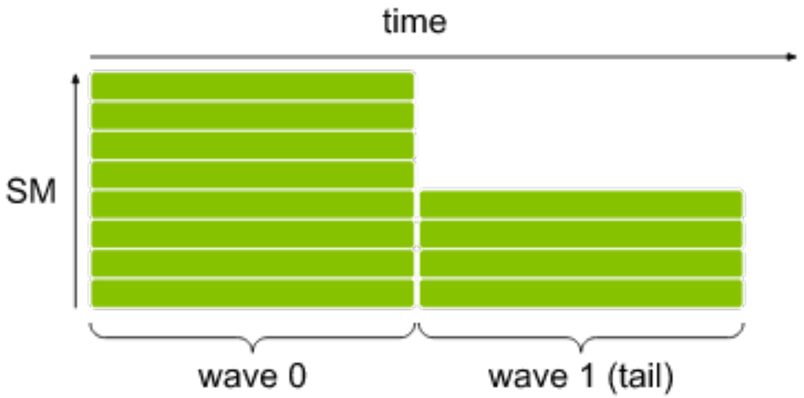

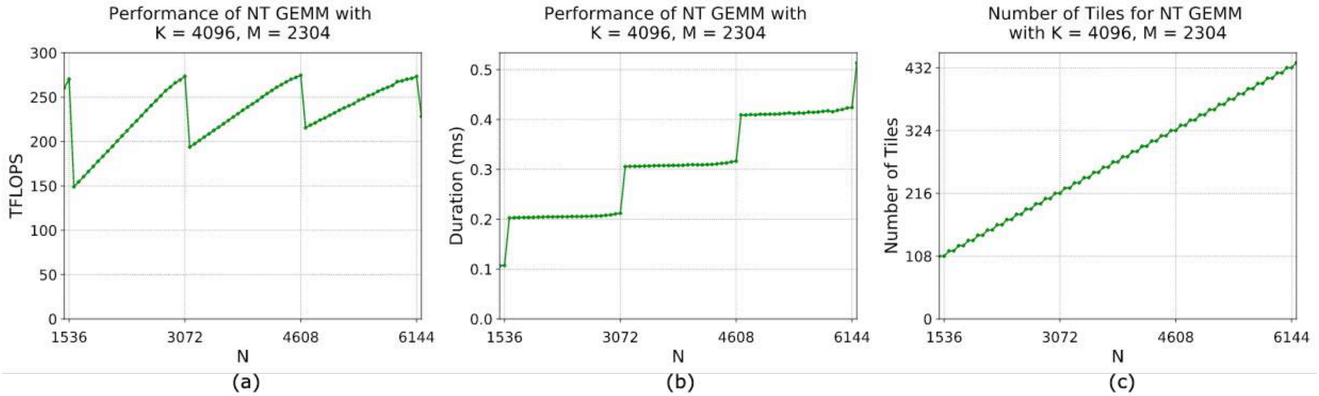

docs.nvidia.com/deeplearning/performance/dl-performance-matrix-multiplication/index.html

Memory access

- GPU has a separate memory unit (called device memory)
- Need to copy from host memory and back (PCIe 4.0 x16 — 32GB/s peak)
- Memory transfer is often a bottleneck
- Pinned (page-locked) memory access is much faster
- Memory hierarchy is a thing, just like on CPUs!

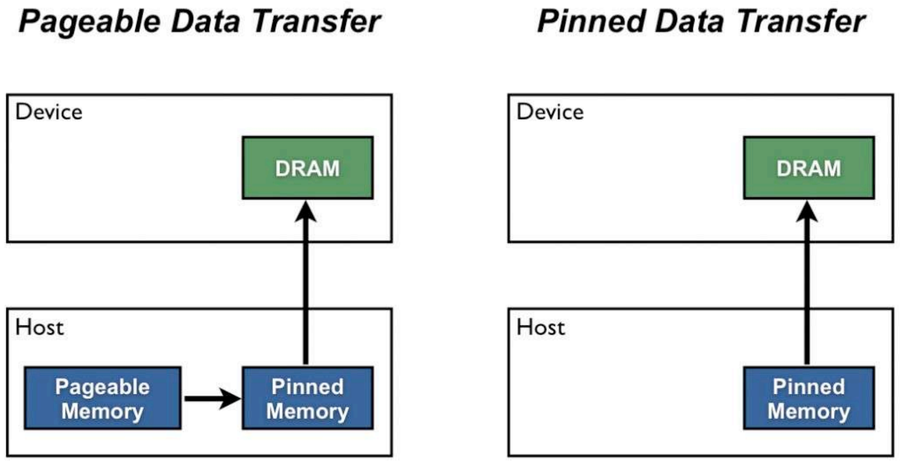

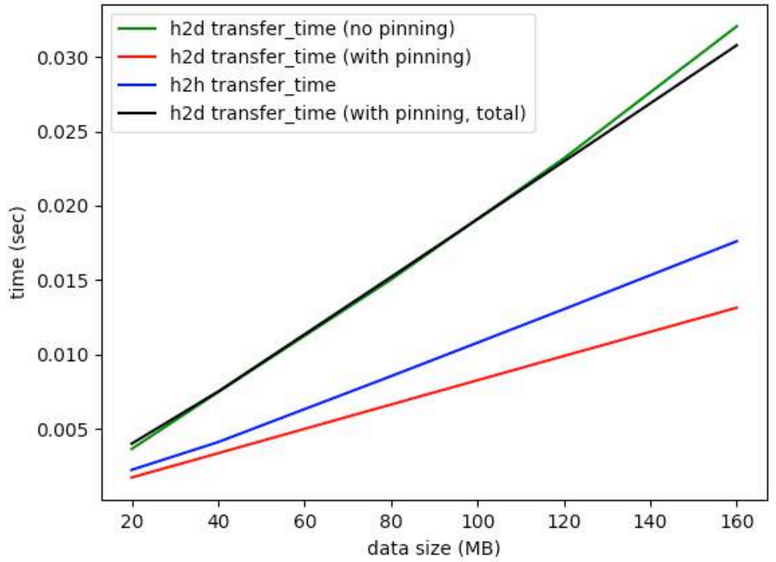

https://developer.nvidia.com/blog/how-optimize-data-transfers-cuda-cc/

Memory access

- GPU has a separate memory unit (called device memory)
- Need to copy from host memory and back (PCIe 4.0 x16 — 32GB/s peak)
- Memory transfer is often a bottleneck
- Pinned (page-locked) memory access is much faster
- Memory hierarchy is a thing, just like on CPUs!

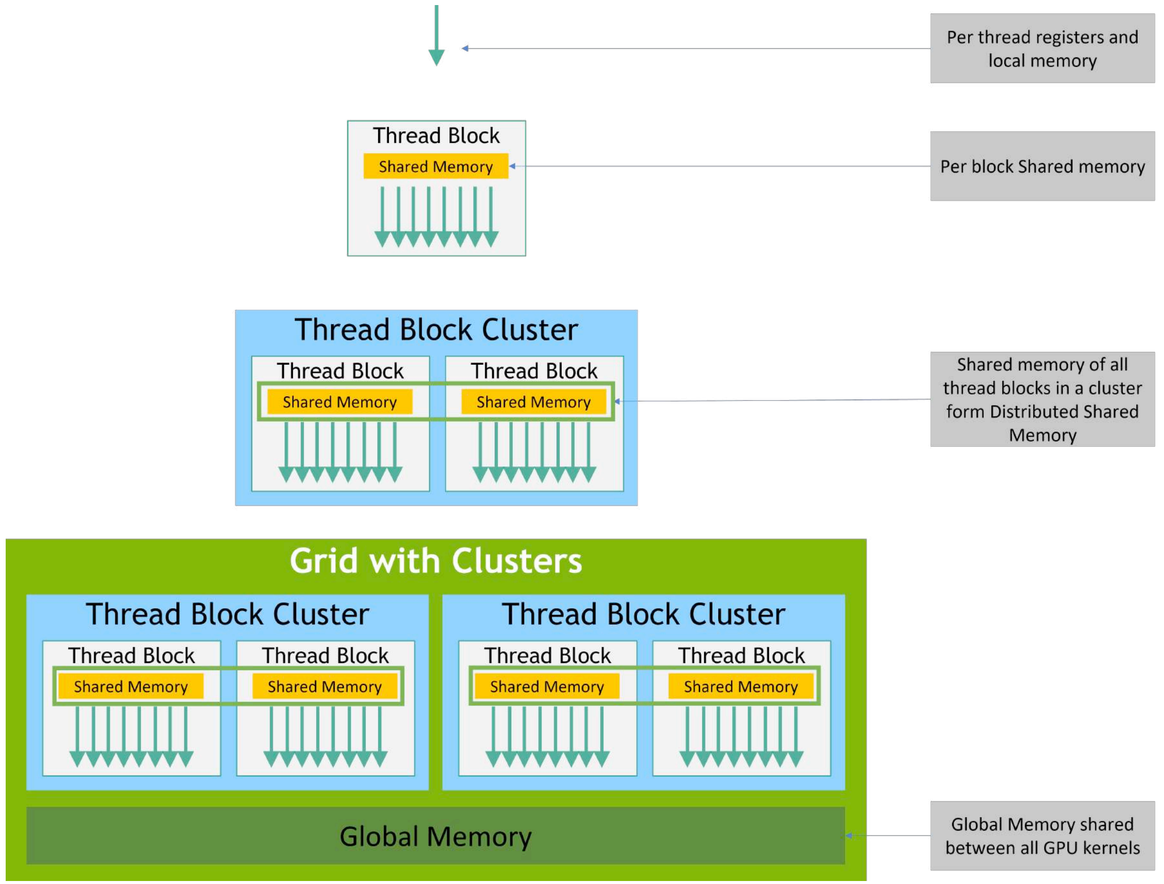

https://developer.nvidia.com/blog/how-optimize-data-transfers-cuda-cc/

Asynchronous execution

- By default, CUDA kernel calls and device transfers are asynchronous
- You can send several kernels and wait for results
- Latest versions of CUDA offer better concurrency mechanisms (streams, graphs)

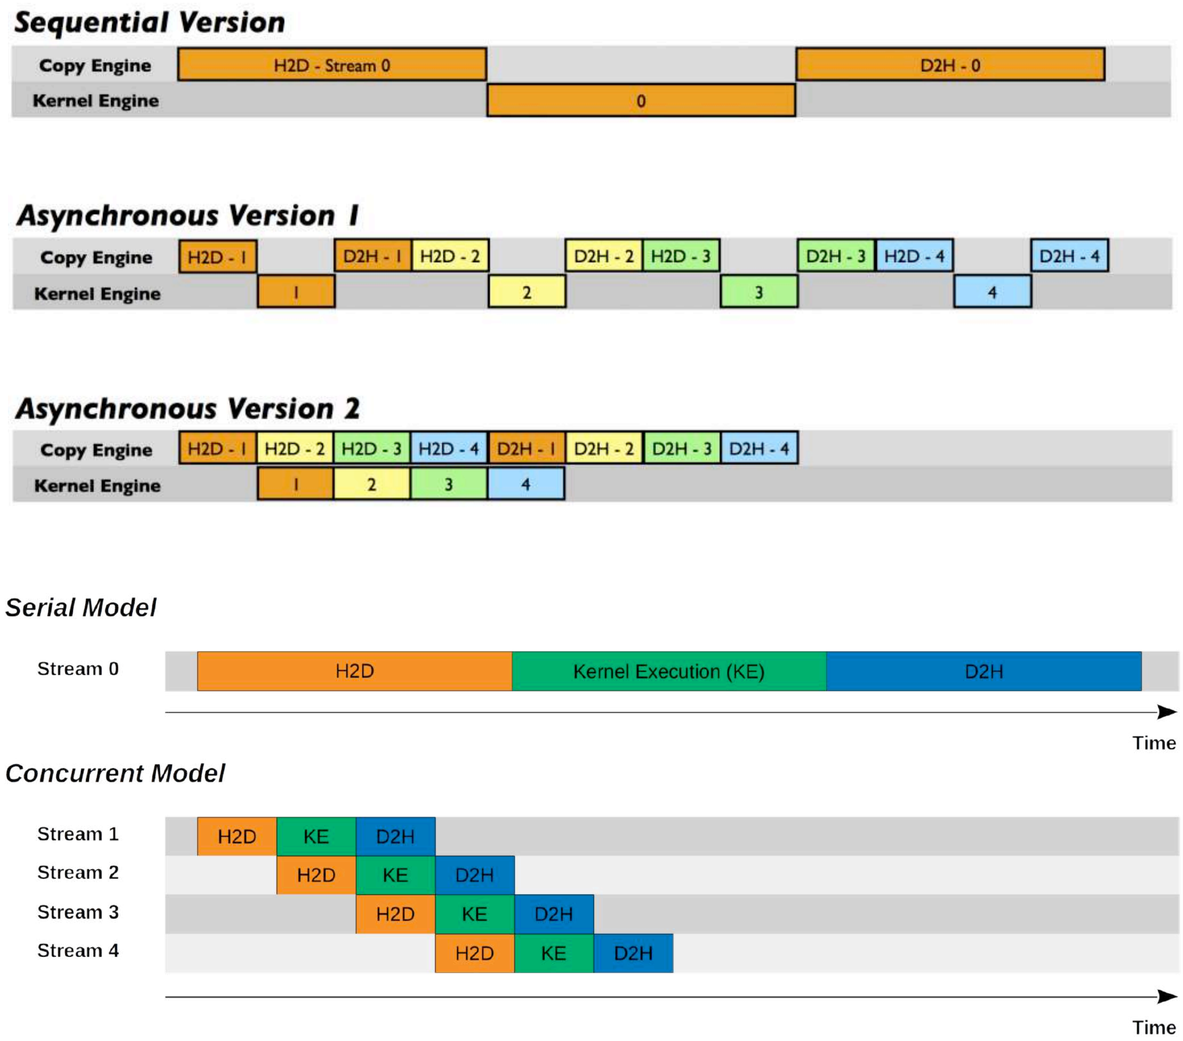

https://developer.nvidia.com/blog/how-overlap-data-transfers-cuda-cc/

DL specifics

With PyTorch as an example:

- Kernel execution is asynchronous, which hides the latency of Python
- Be careful when benchmarking though!
- Calling Tensor.item() triggers a D2H copy
- Allocated memory is not released immediately to simplify caching
- torch.backends.cudnn.benchmark=True
- CUDA streams, graphs etc. are available in latest releases

nn.Conv2d with 64 3x3 filters applied to an input with batch size = 32, channels = width = height = 64.

Setting	`cudnn.benchmark = False` (the default)	`cudnn.benchmark = True`	Speedup
Forward propagation (FP32) [us]	1430	840	1.70
Forward + backward propagation (FP32) [us]	2870	2260	1.27

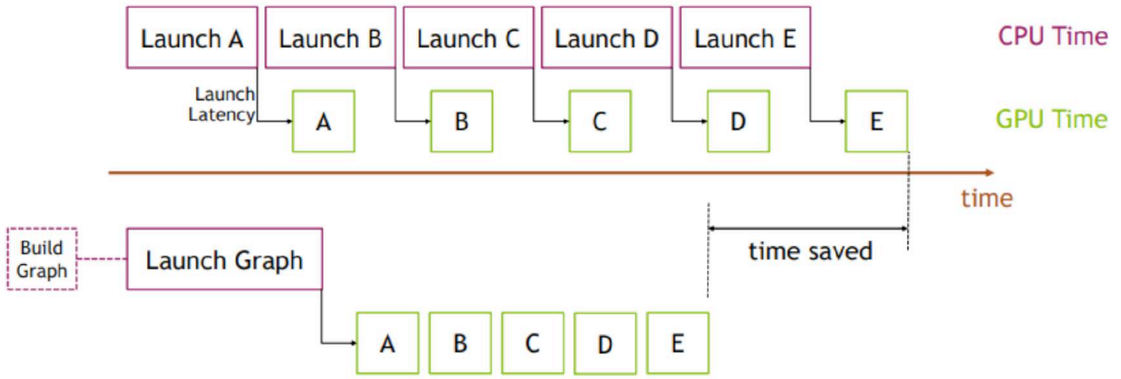

Measuring performance

- Benchmarking is a key step of understanding your bottlenecks and measuring the impact of optimizations
- Basically, just run the code several times or measure large workloads
- Can be done via %timeit or timeit.Timer (mind the synchronization)
- Due to possible side-effects (preallocation, caching), warmup and randomization are often necessary
- In PyTorch, you can use torch.utils.benchmark
- **Don't overoptimize!**