

MLOps, k8s, GitOps and other acronyms

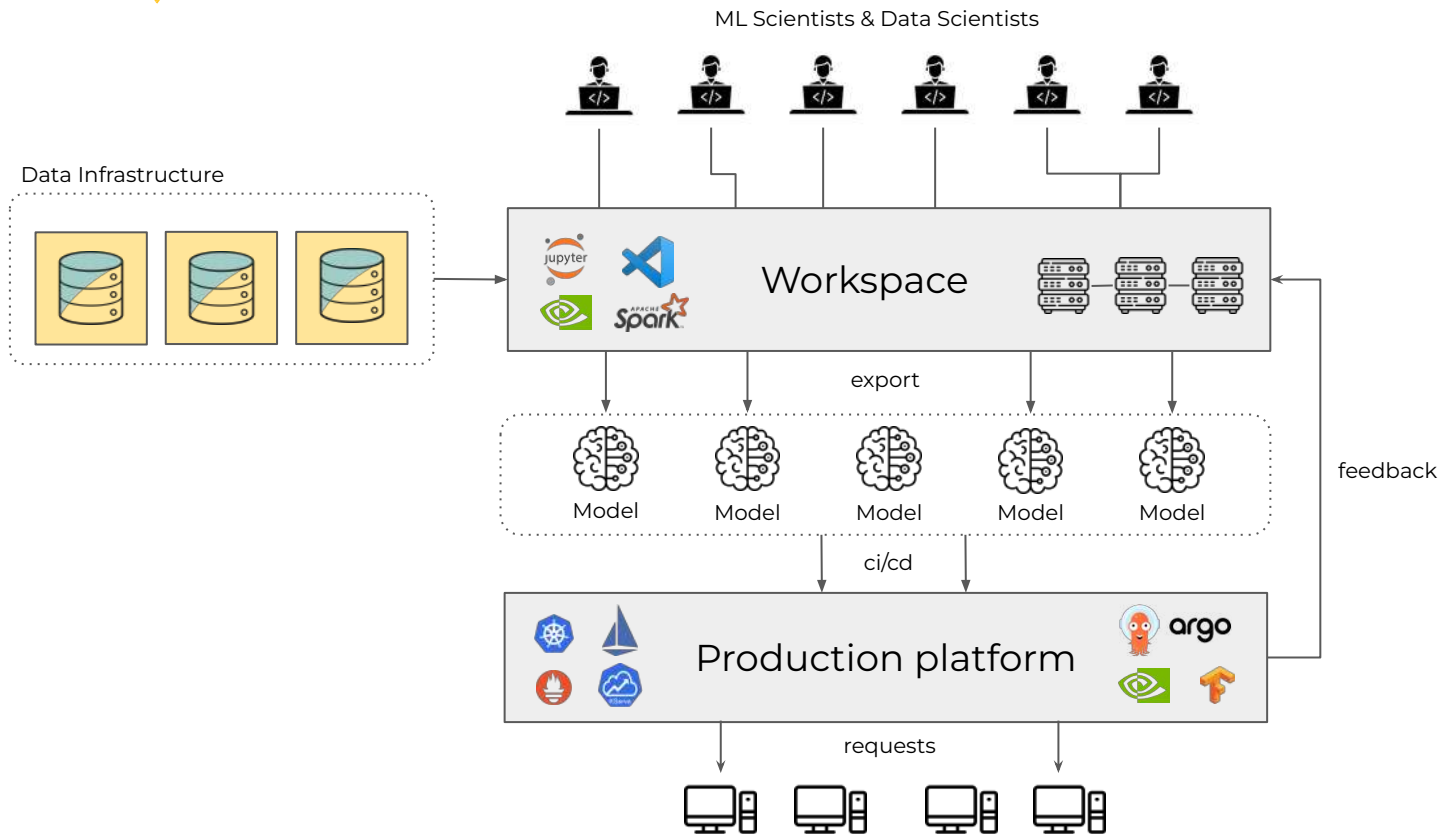
Automating multi-cluster kubernetes environments for ML tasks and services




Gleb Vazhenin

- Born in Cherepovets, Russia
- Bachelor of science in Nuclear Physics (NRNU MEPHI)
- 8 Years of DS/MLE/SRE/DevOps experience
- Staff MLE @ Bumble
- Wish to become chess grandmaster

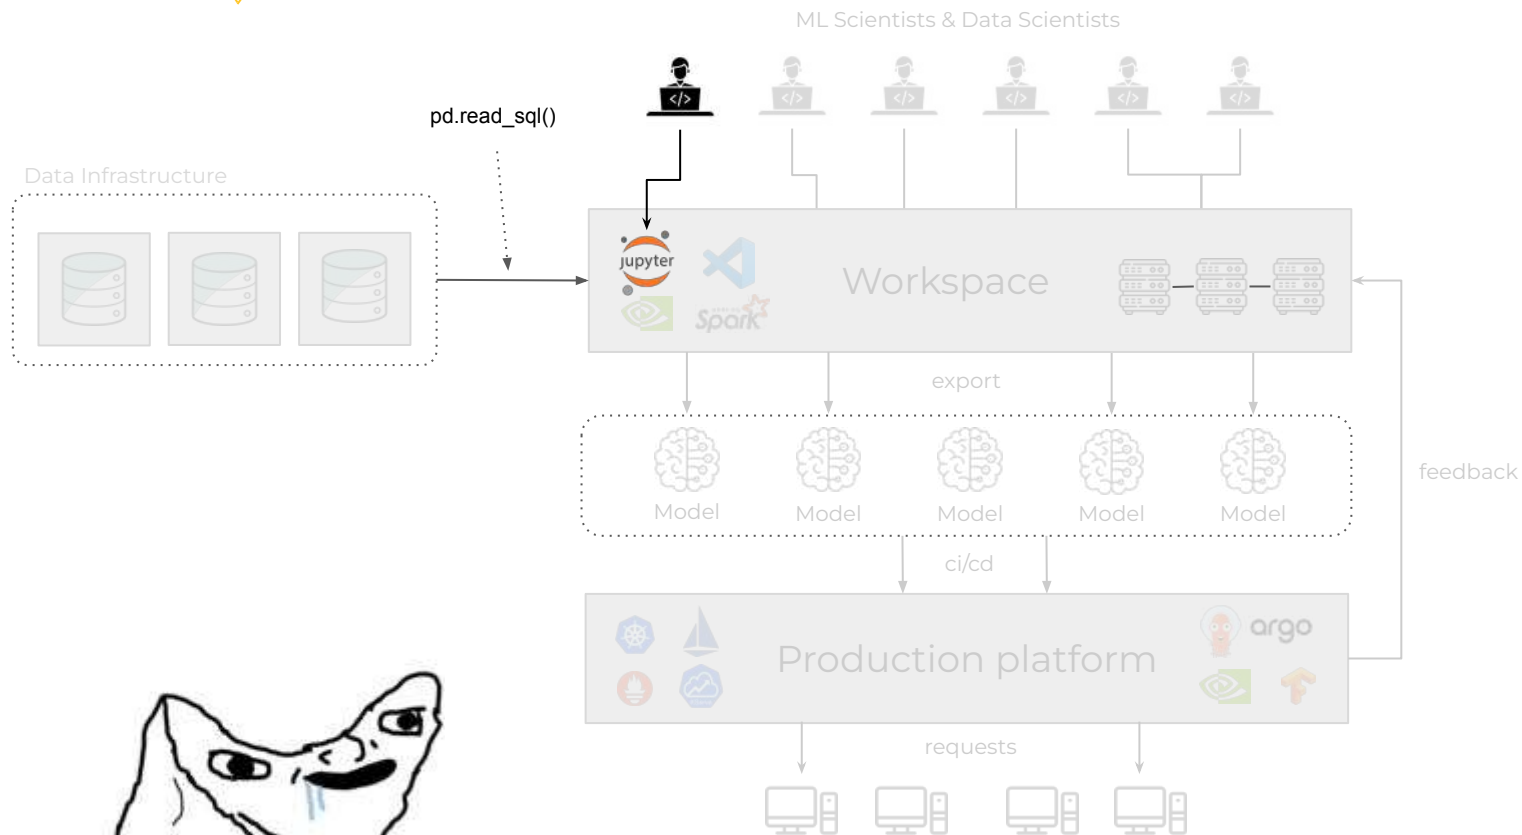
MLOps



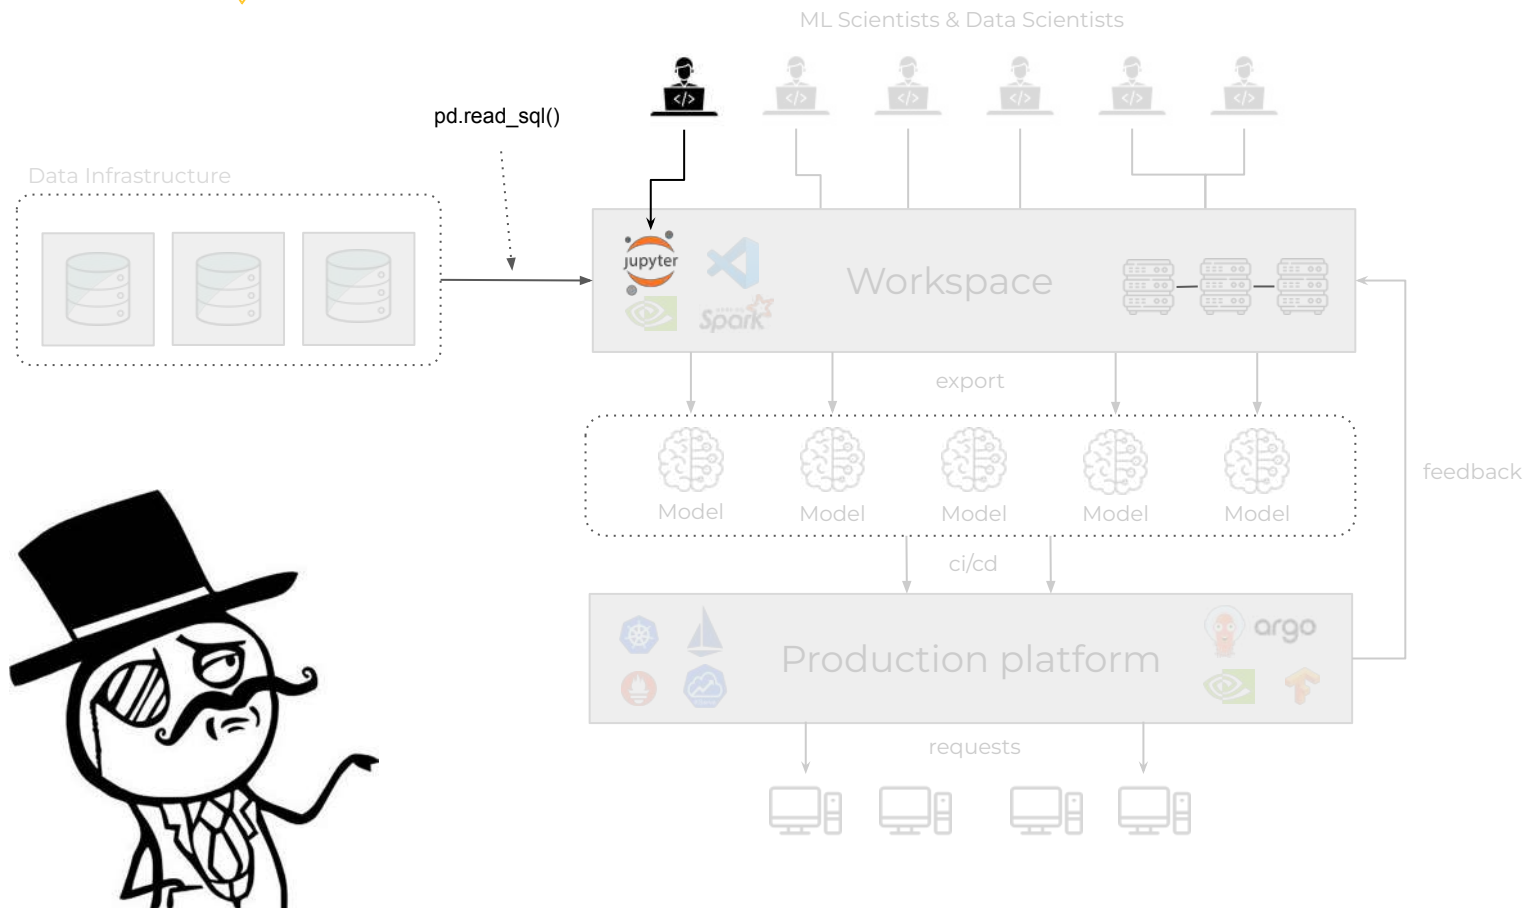
MLOps?



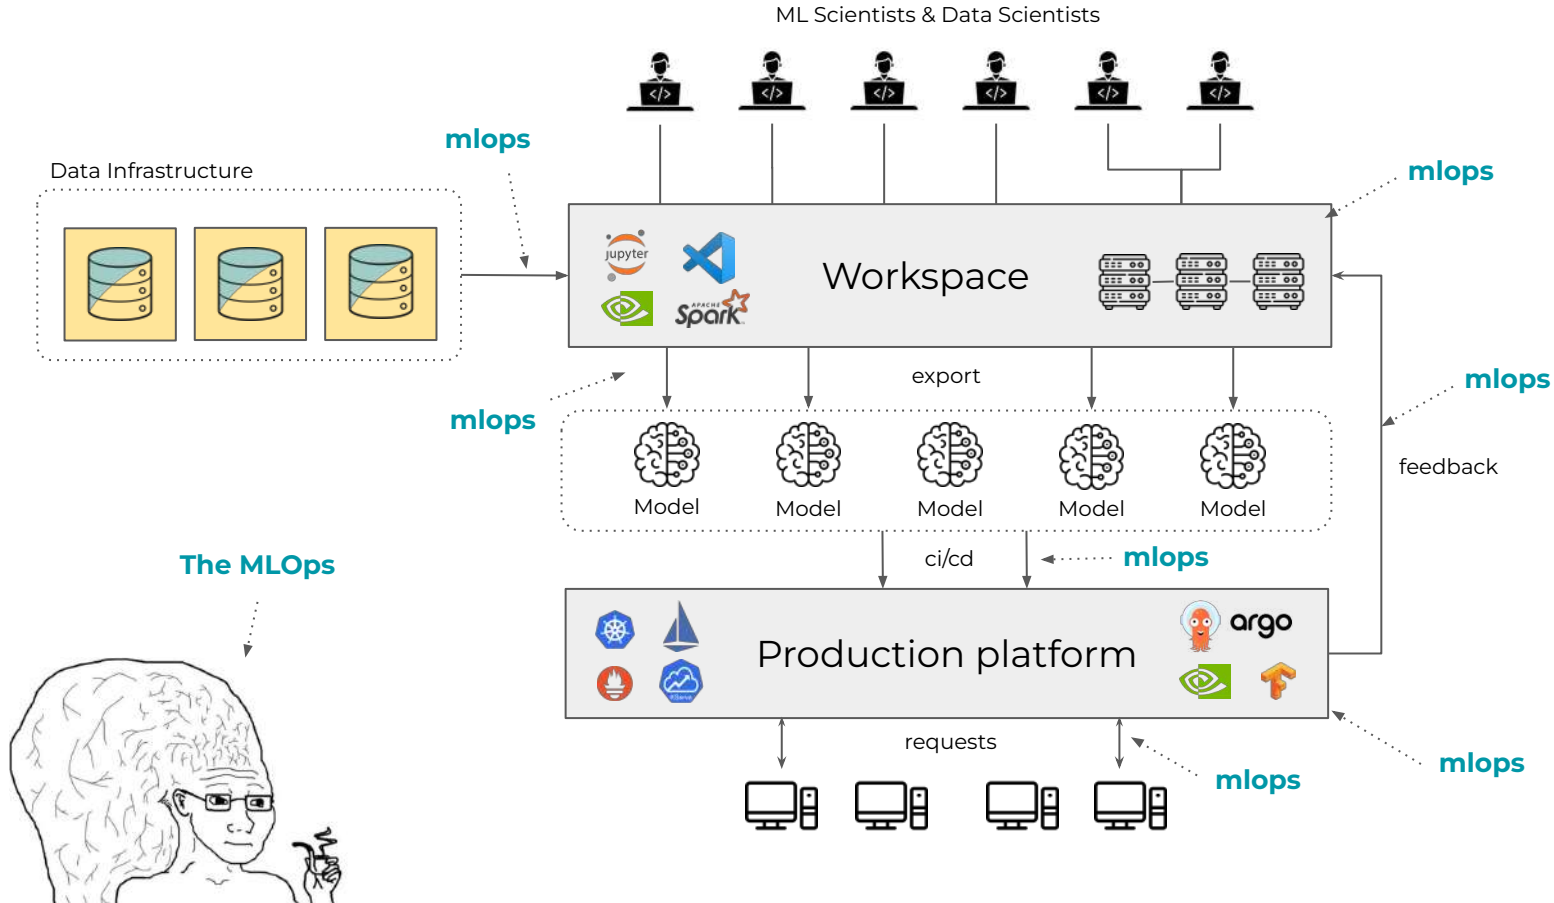
MLOps?



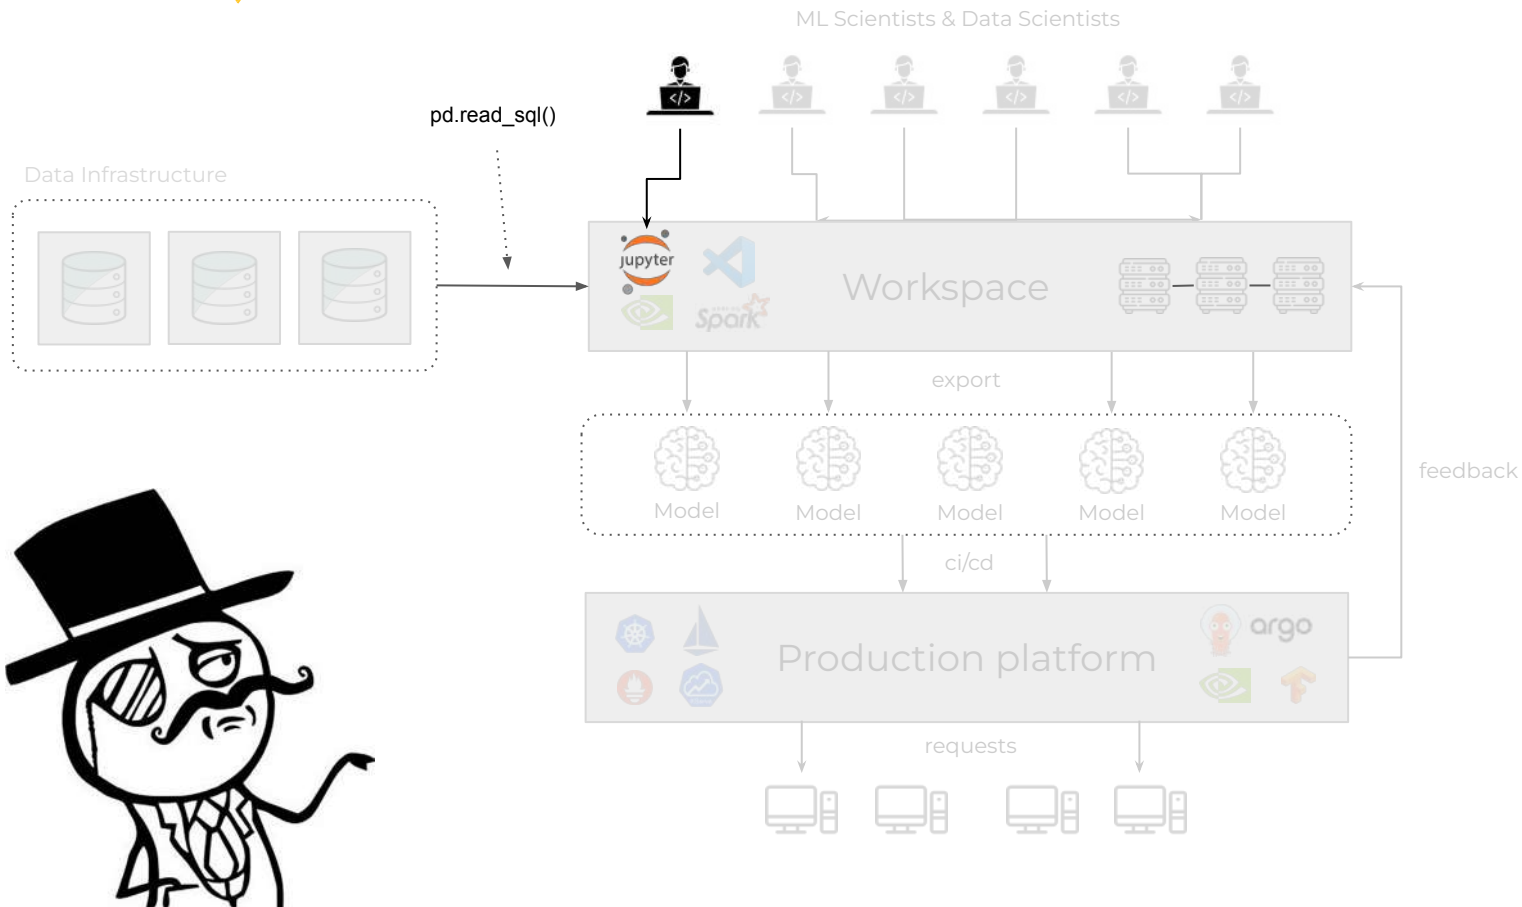
MLOps?



MLOps - is the core function of the Data Science team



That's the desired way!



**What if there are so many different services and tools,
that you don't want to hire a specialist to maintain each
one?**



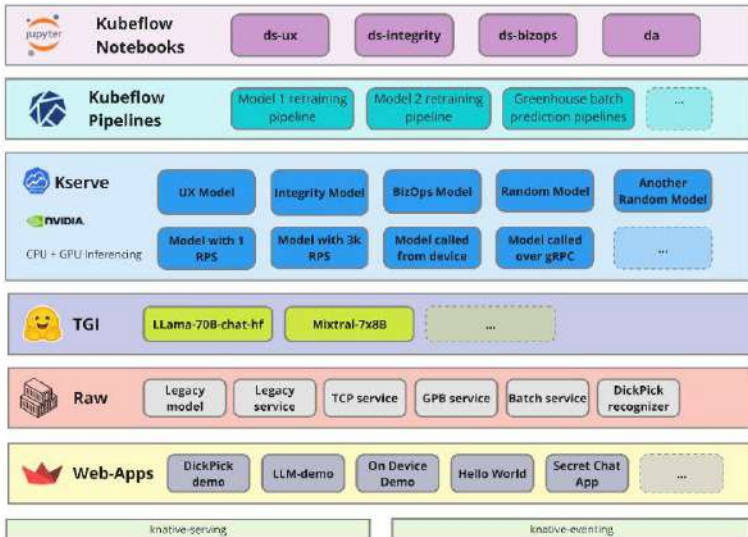


Kubernetes cluster (Prod - Europe)

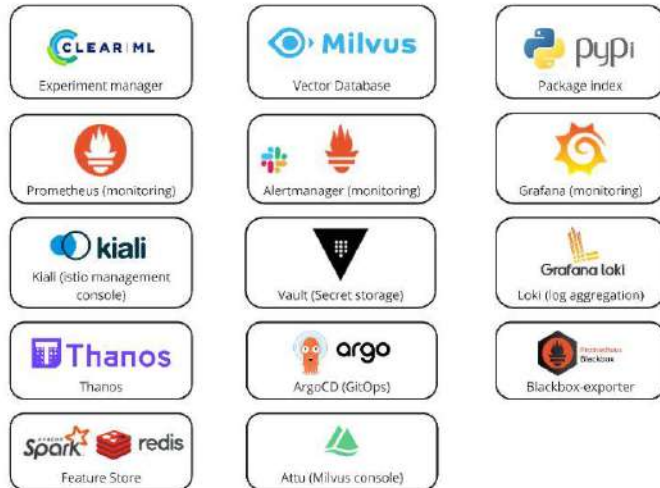
istio-ingress-gateway

ds-bumble-dev-gateway

ML Models and Workspace



Services



Kubernetes control plane

Control plane VMs

- 5x16 GB Memory
- control-plane
- 5x4 CPU
- etcd

metalib gpu-device-plugin



vector

metalib gpu-device-plugin



vector

nfs-provisioner

cephfs-provisioner



The K8s

Control Plane (master nodes)



Controller
Manager



APIServer



etcd



Scheduler

manifest sample

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: private-detector
  labels:
    app: pd
spec:
  replicas: 3
  template:
    metadata:
      labels:
        app: pd
    spec:
      containers:
        - name: pd
          image: pd:1.48.8
          ports:
            - containerPort: 80
```

Worker 1

kubelet

kube-proxy



Worker 2

kubelet

kube-proxy



Worker 3

kubelet

kube-proxy



Why K8s?

- **Scalability**
- Fault Tolerance and High Availability
- Declarative Configuration and Automation
- Portability
- Ecosystem and Community
- Containers
- Rolling upgrades and rollbacks
- Service discovery and load balancing

You can't have scalability problems if nobody uses your app




Why K8s?

- Scalability
- **Fault Tolerance and High Availability**
- Declarative Configuration and Automation
- Portability
- Ecosystem and Community
- Containers
- Rolling upgrades and rollbacks
- Service discovery and load balancing



Why K8s?

- Scalability
- Fault Tolerance and High Availability
- **Declarative Configuration and Automation**
- Portability
- Ecosystem and Community
- Containers
- Rolling upgrades and rollbacks
- Service discovery and load balancing




Why K8s?

- Scalability
- Fault Tolerance and High Availability
- Declarative Configuration and Automation
- **Portability**
- Ecosystem and Community
- Containers
- Rolling upgrades and rollbacks
- Service discovery and load balancing



Why K8s?

- Scalability
- Fault Tolerance and High Availability
- Declarative Configuration and Automation
- Portability
- **Ecosystem and Community**
- Containers
- Rolling upgrades and rollbacks
- Service discovery and load balancing



Why K8s?

- Scalability
- Fault Tolerance and High Availability
- Declarative Configuration and Automation
- Portability
- Ecosystem and Community
- **Containers**
- Rolling upgrades and rollbacks
- Service discovery and load balancing




Why K8s?

- Scalability
- Fault Tolerance and High Availability
- Declarative Configuration and Automation
- Portability
- Ecosystem and Community
- Containers
- **Rolling upgrades and rollbacks**
- Service discovery and load balancing

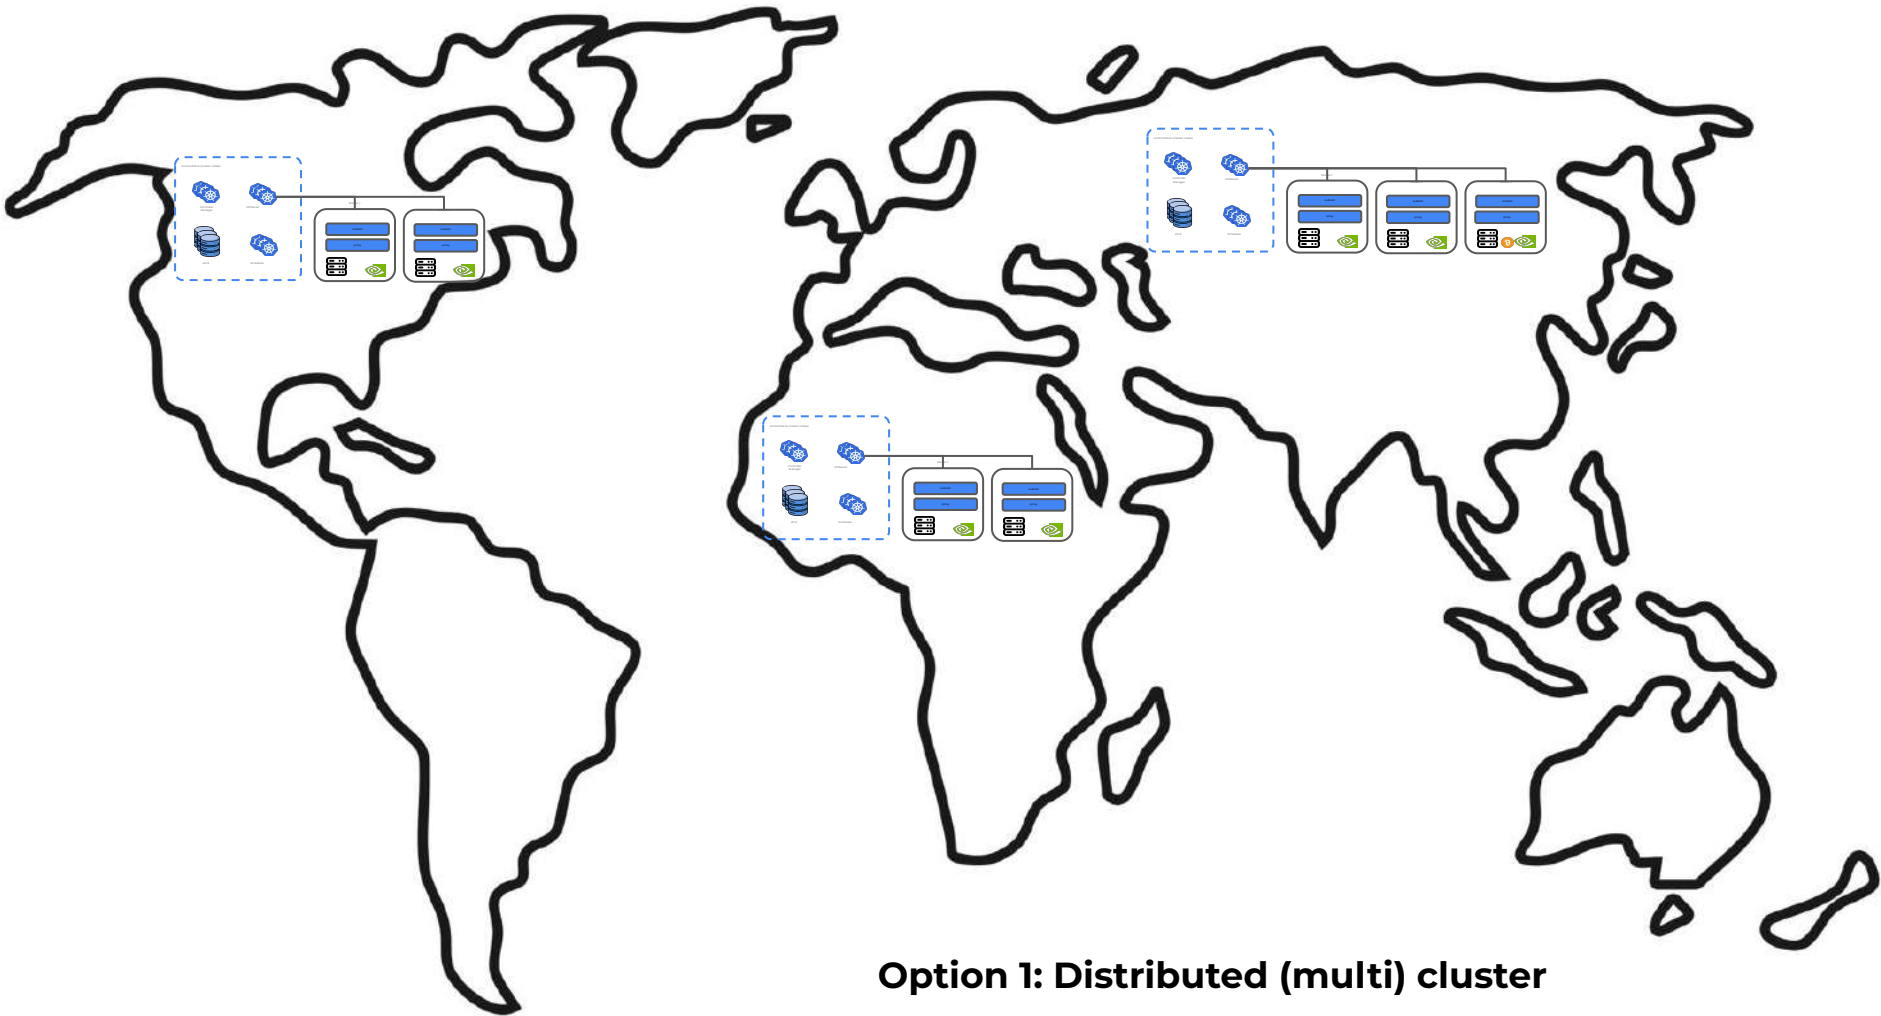
Why K8s?

- Scalability
- Fault Tolerance and High Availability
- Declarative Configuration and Automation
- Portability
- Ecosystem and Community
- Containers
- Rolling upgrades and rollbacks
- **Service discovery and load balancing**

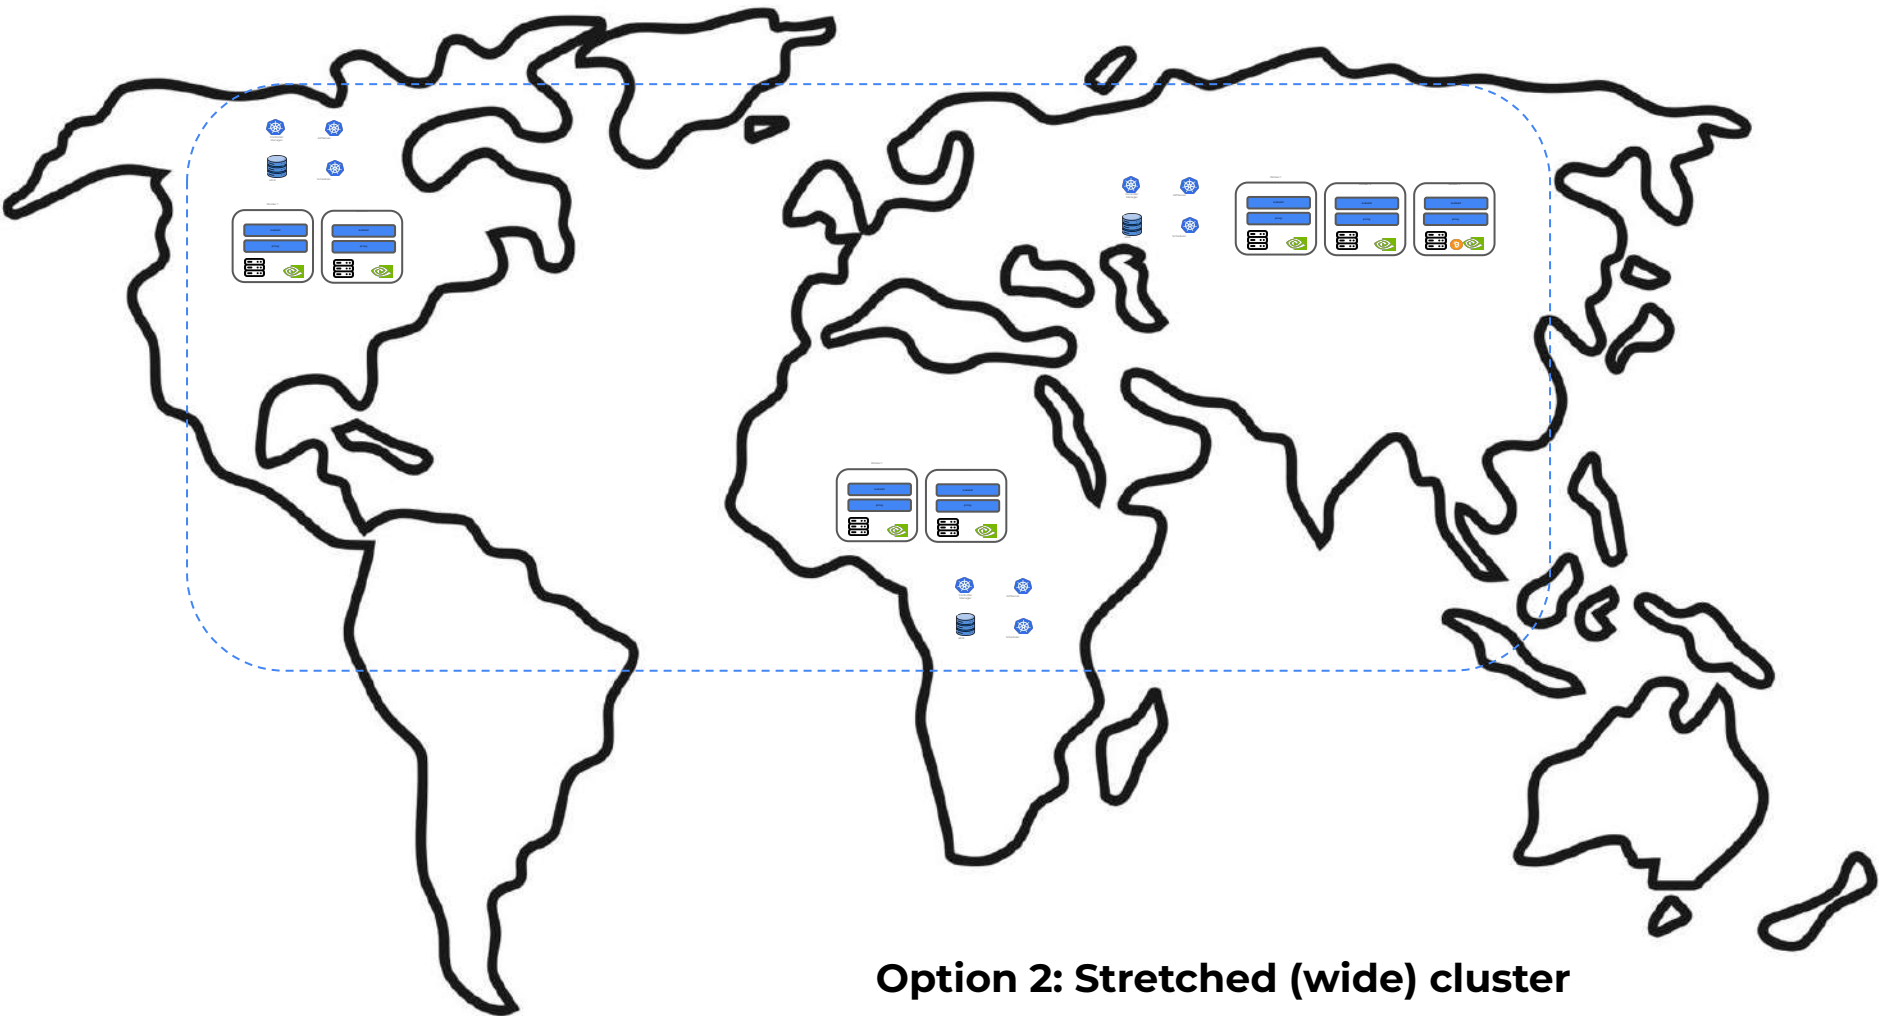


What if the DC is so cheap, that there's a disaster in electricity every other week?






Option 1: Distributed (multi) cluster



Option 2: Stretched (wide) cluster

Wide vs Multi - Key points

- **Redundancy** - same
- **Scalability** - same
- **Latency**
 - Wide - **higher** latencies
 - Multi - **lower** latencies
- **Isolation**
 - Wide - same cluster, **less** isolation
 - Multi - different clusters, **strong** isolation by design
- **Management complexity**
 - Wide - **easier** to maintain (same cluster), though not out-of-the-box
 - Multi - **harder** to maintain

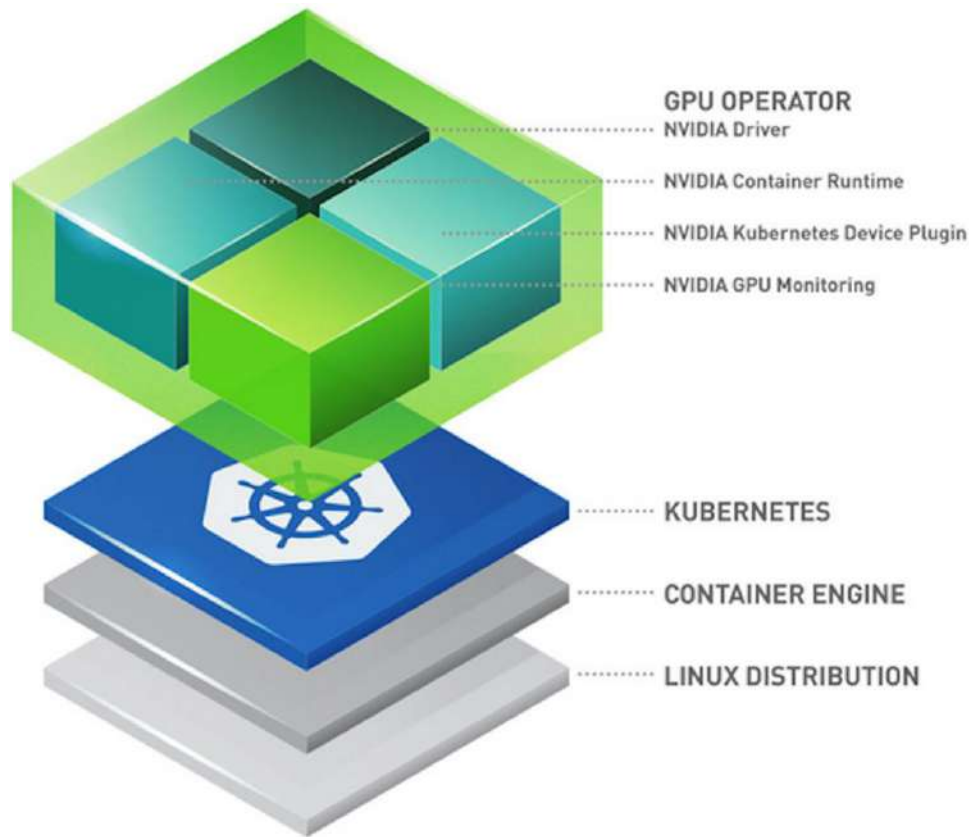


GPUs on k8s

NVIDIA GPU Operator

Components:

- GPU Feature discovery
- Nvidia Container Runtime
- K8s Device Plugin
- DCGM Exporter
- Driver Manager
- MIG Manager



NVIDIA GPU Operator

Components:

- **GPU Feature discovery**
- Nvidia Container Runtime
- K8s Device Plugin
- DCGM Exporter
- Driver Manager
- MIG Manager

NVIDIA/gpu-feature-discovery

GPU plugin to the node feature discovery for Kubernetes



17

Contributors

13

Issues

255

Stars

44

Forks



```
$ kubectl get nodes --show-labels
```

NAME	STATUS	ROLES	AGE	VERSION	LABELS
ds-node1	Ready	worker	210d	v1.24.3	...,nvidia.com/gpu.product=NVIDIA-A100,nvidia.com/gpu.replicas=2
ds-node2	Ready	worker	210d	v1.24.3	...,nvidia.com/gpu.product=Tesla-T4,nvidia.com/gpu.replicas=4
ds-node3	Ready	worker	210d	v1.24.3	...,nvidia.com/gpu.product=Tesla-T4,nvidia.com/gpu.replicas=4

NVIDIA GPU Operator

Components:

- GPU Feature discovery
- **Nvidia Container Runtime**
- K8s Device Plugin
- DCGM Exporter
- Driver Manager
- MIG Manager

GPU-Accelerated
Applications



Container
Technologies



cri-o ...

NVIDIA CONTAINER RUNTIME



NVIDIA GPU Operator

Components:

- GPU Feature discovery
- Nvidia Container Runtime
- **K8s Device Plugin**
- DCGM Exporter
- Driver Manager
- MIG Manager

NVIDIA/k8s-
device-plugin

NVIDIA device plugin for Kubernetes



28
Contributors

9
Used by

2k
Stars

556
Forks



```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: private-detector
  labels:
    app: pd
spec:
  replicas: 3
  template:
    metadata:
      labels:
        app: pd
    spec:
      containers:
        - name: pd
          image: pd:1.48.8
          ports:
            - containerPort: 80
          resources:
            limits:
              cpu: 2
              memory: 16Gi
              nvidia.com/gpu: 2
```

NVIDIA GPU Operator

Components:

- GPU Feature discovery
- Nvidia Container Runtime
- K8s Device Plugin
- **DCGM Exporter**
- Driver Manager
- MIG Manager



NVIDIA GPU Operator

Components:

- GPU Feature discovery
- Nvidia Container Runtime
- K8s Device Plugin
- DCGM Exporter
- **Driver Manager**
- MIG Manager

Fri Dec 30 10:57:52 2022

```

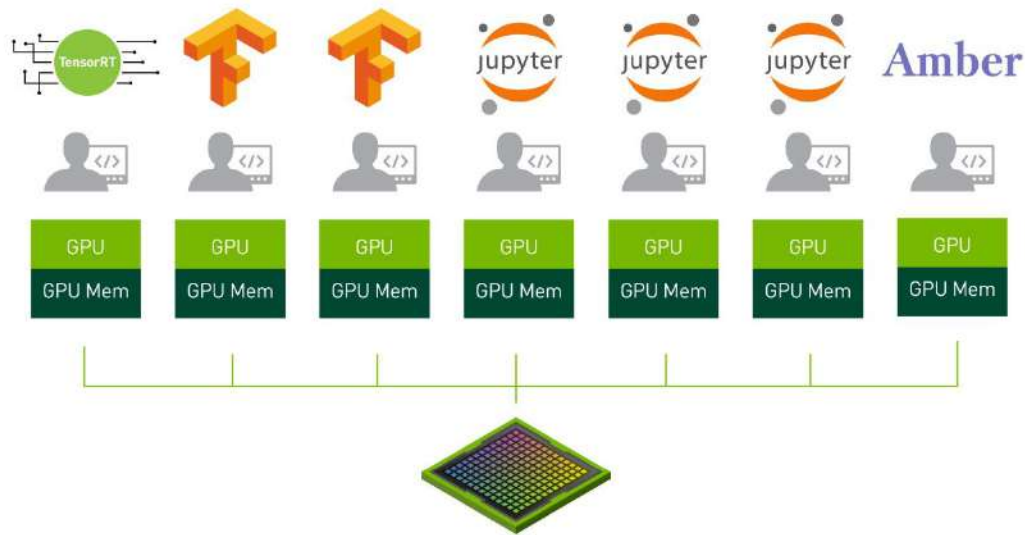
+-----+
| NVIDIA-SMI 525.60      Driver Version: 525.60      CUDA Version: 12.0      |
+-----+-----+-----+
| GPU   Name              Persistence-M| Bus-Id        Disp.A | Volatile Uncorr. ECC |
| Fan   Temp   Perf   Pwr:Usage/Cap|      Memory-Usage | GPU-Util  Compute M. |
|=====+=====+=====+
|    0   NVIDIA A100-PCI...    Off   | 00000000:86:00.0 Off  |                     0 |
| N/A    30C      P0      36W / 250W |      0MiB / 40960MiB |           0%      Default |
|                                     |                     |             Disabled |
+-----+-----+-----+
| Processes:                                                       |
| GPU   GI    CI          PID    Type    Process name                      GPU Memory |
|          ID    ID                                   Usage                      |
|=====+=====+=====+
| No running processes found                                         |
+-----+

```

NVIDIA GPU Operator

Components:

- GPU Feature discovery
- Nvidia Container Runtime
- K8s Device Plugin
- DCGM Exporter
- Driver Manager
- **MIG Manager**



NVIDIA GPU Operator

Components:

- GPU Feature discovery
- Nvidia Container Runtime
- K8s Device Plugin
- DCGM Exporter
- Driver Manager
- **MIG Manager**

```
kubectl label nodes ds-node1 nvidia.com/mig.config=all-1g.10gb
```

```
kubectl label nodes ds-node1 nvidia.com/mig.config=all-1g.5gb
```

```
kubectl label nodes ds-node1 nvidia.com/mig.config=all-3g.40gb
```

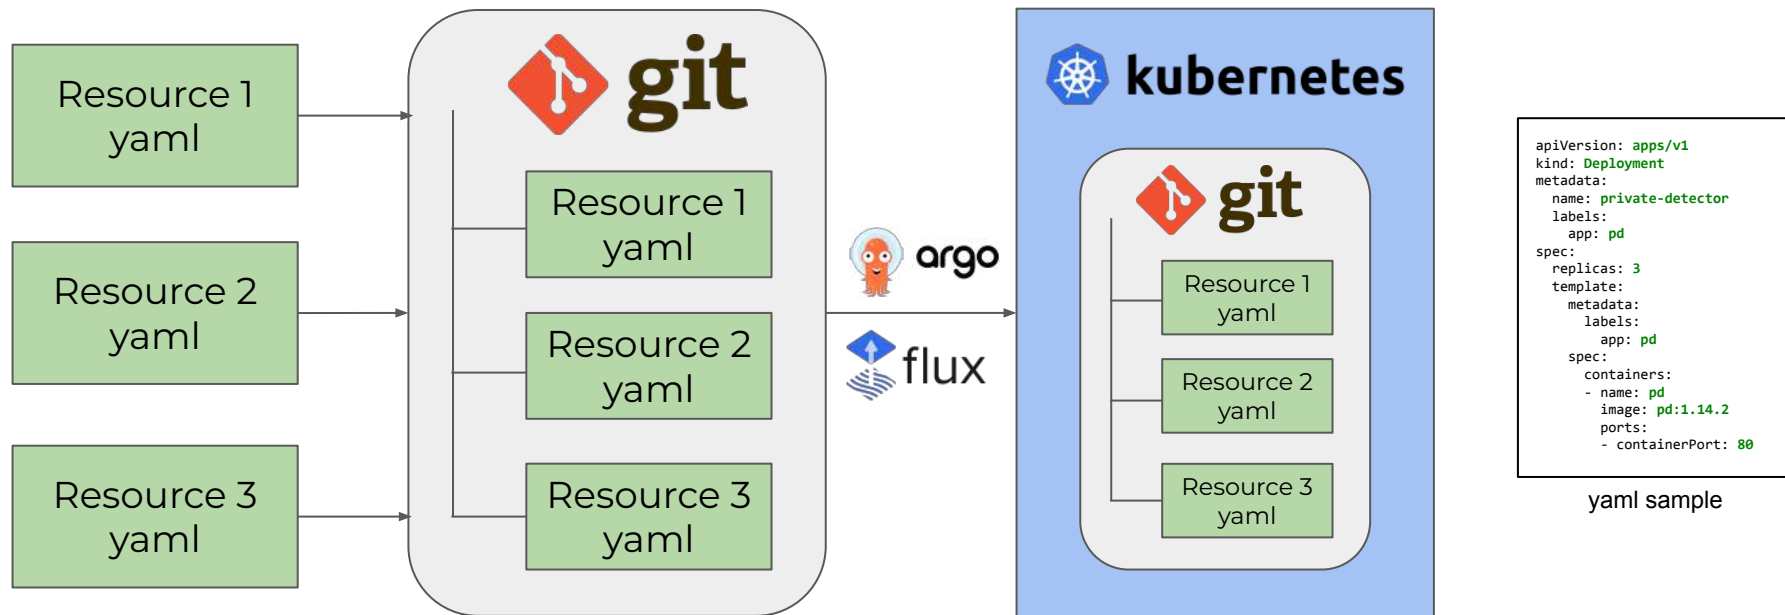
```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: private-detector
  labels:
    app: pd
spec:
  replicas: 3
  template:
    metadata:
      labels:
        app: pd
    spec:
      containers:
        - name: pd
          image: pd:1.76.9
          ports:
            - containerPort: 80
          resources:
            limits:
              cpu: 1
              memory: 2Gi
              nvidia.com/mig-1g.5gb: 1
```


What if I need to manage 50 GPU clusters simultaneously, taking into account different service configurations?

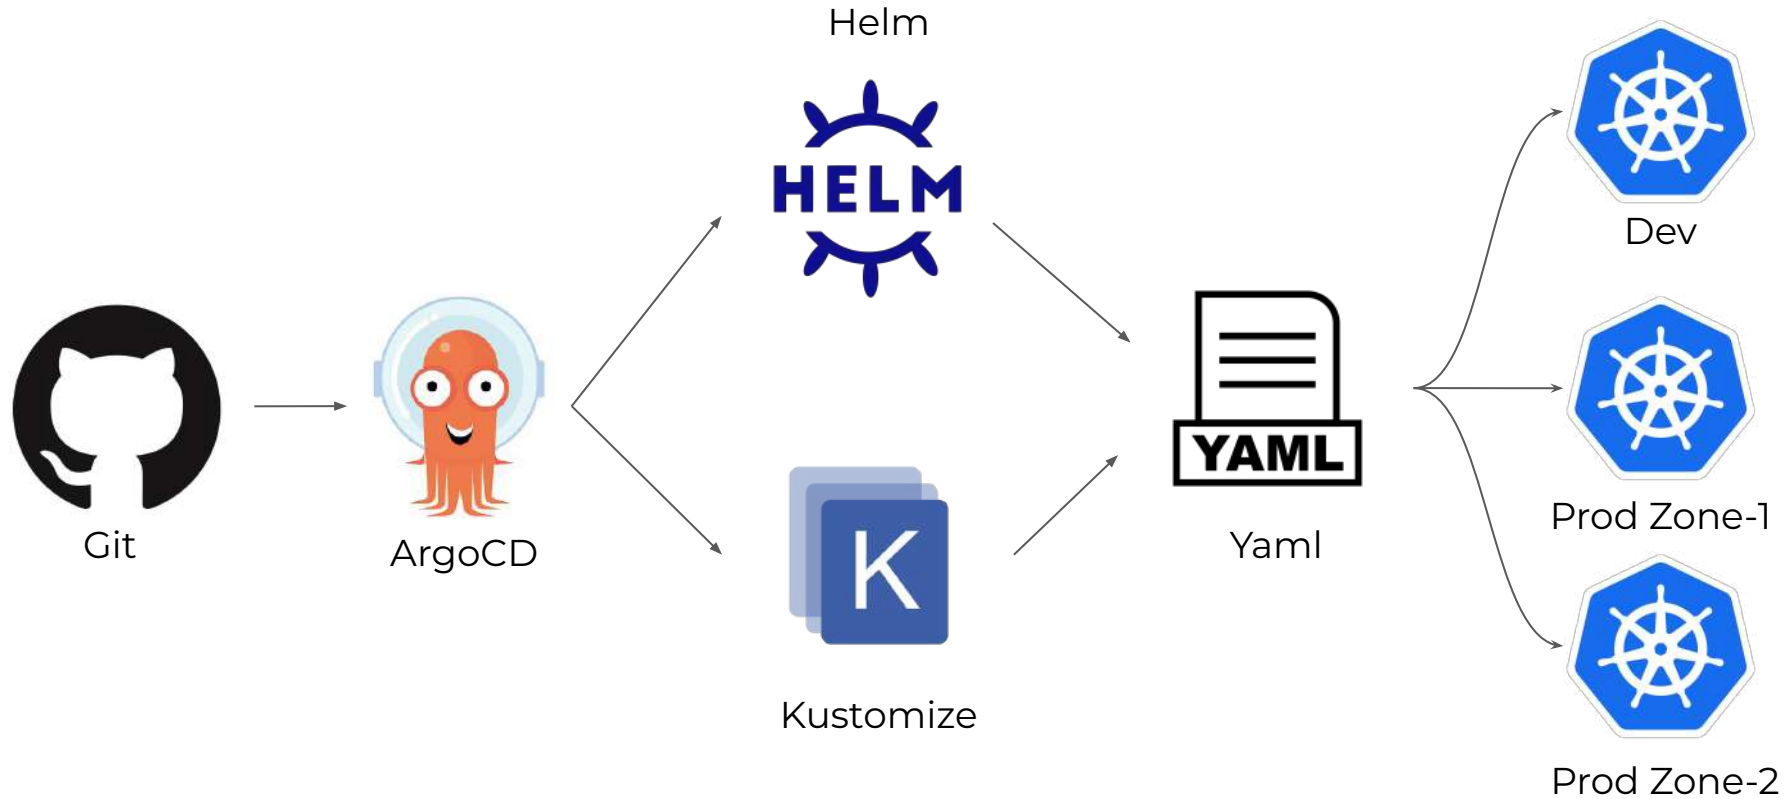


GitOps

GitOps modernises software management by allowing Engineers to declaratively manage infrastructure and software code using a single source of truth — typically a Git repository.



Kustomize & Helm are the tools used to manage (template) kubernetes manifests



Smooth GitOps experience: Overlays



apps

private-detector

base

deployment.yaml

bumble

deployment_patch.yaml

kustomization.yaml

zone-1-dev

deployment_zone_1_dev_patch.yaml

kustomization.yaml

zone-1-prod

deployment_zone_1_prod_patch.yaml

kustomization.yaml

zone-2-prod

deployment_zone_2_patch.yaml

kustomization.yaml

apps/private-detector/base/deployment.yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: private-detector
  labels:
    app: pd
spec:
  replicas: 3
  template:
    metadata:
      labels:
        app: pd
    spec:
      containers:
        - name: pd
          image: pd:1.14.2
          ports:
            - containerPort: 80
```

Shipped like this by private-detector community

apps/private-detector/bumble/deployment_patch.yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: private-detector
spec:
  template:
    spec:
      image: pd:1.13.2
```

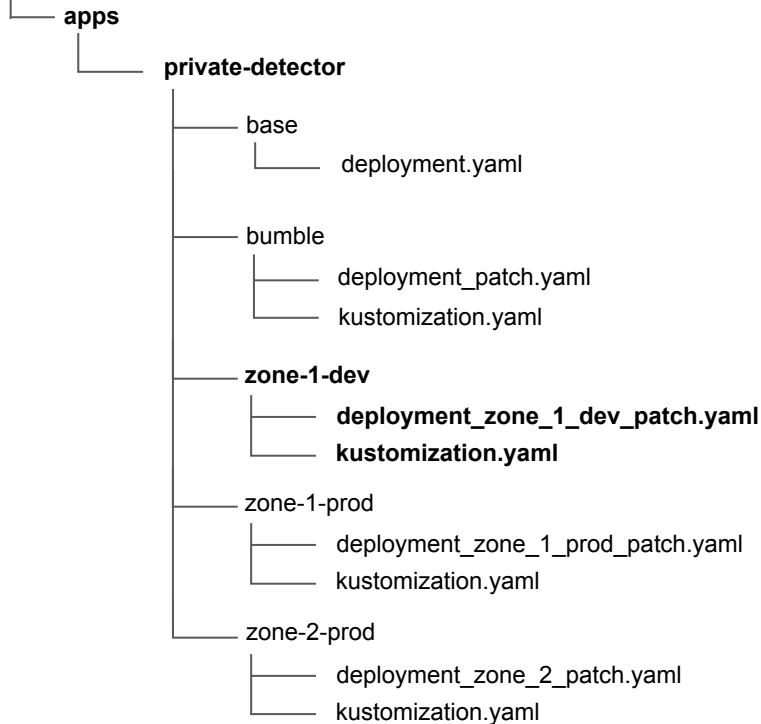
apps/private-detector/bumble/kustomization.yaml

```
resources:
- ../base/deployment.yaml

patchesStrategicMerge:
- deployment_patch.yaml
```

Bumble cluster-agnostic label patch

Smooth GitOps experience: Overlays



apps/private-detector/**zone-1-dev**/deployment_zone_1_dev_patch.yaml

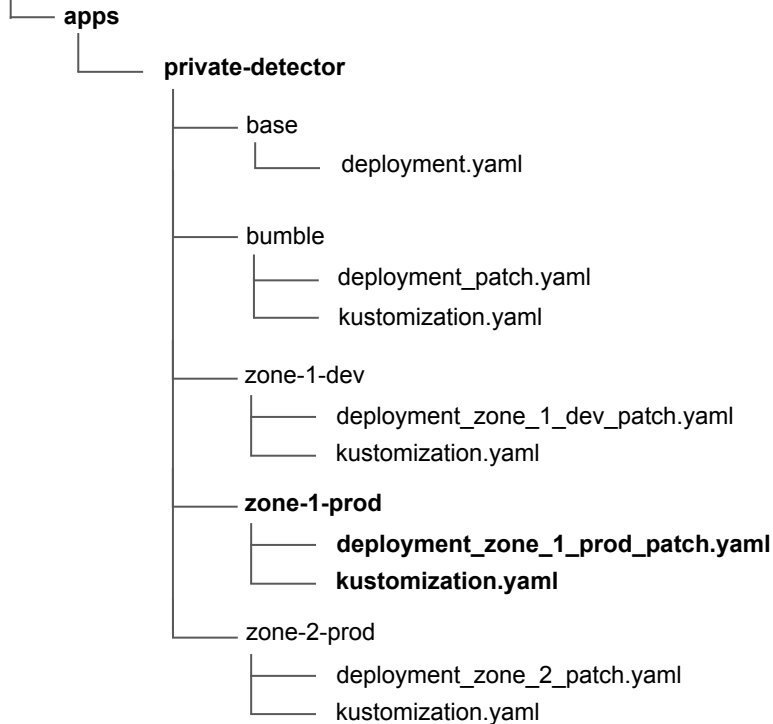
```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: private-detector
spec:
  replicas: 1
  template:
    spec:
      containers:
        - name: pd
          resources:
            limits:
              nvidia.com/mig-1g.5gb: 1
```

apps/private-detector/**zone-1-dev**/kustomization.yaml

```
resources:
- ../bumble

patchesStrategicMerge:
- deployment_zone_1_patch.yaml
```

Smooth GitOps experience: Two-level overlays



apps/private-detector/**zone-1-prod**/deployment_zone_1_prod_patch.yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: private-detector
spec:
  replicas: 10
  template:
    spec:
      nodeSelector:
        nvidia.com/gpu.product: NVIDIA-A100
      containers:
        - name: pd
          resources:
            limits:
              nvidia.com/gpu: 1
```

apps/private-detector/**zone-1-prod**/kustomization.yaml

```
resources:
- ../../../../base/istio/bumble

patchesStrategicMerge:
- deployment_zone_2_patch.yaml
```

Smooth GitOps experience: Two-level overlays

base manifest

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: private-detector
  labels:
    app: pd
spec:
  replicas: 3
  template:
    metadata:
      labels:
        app: pd
    spec:
      containers:
        - name: pd
          image: pd:1.14.2
          ports:
            - containerPort: 80
```

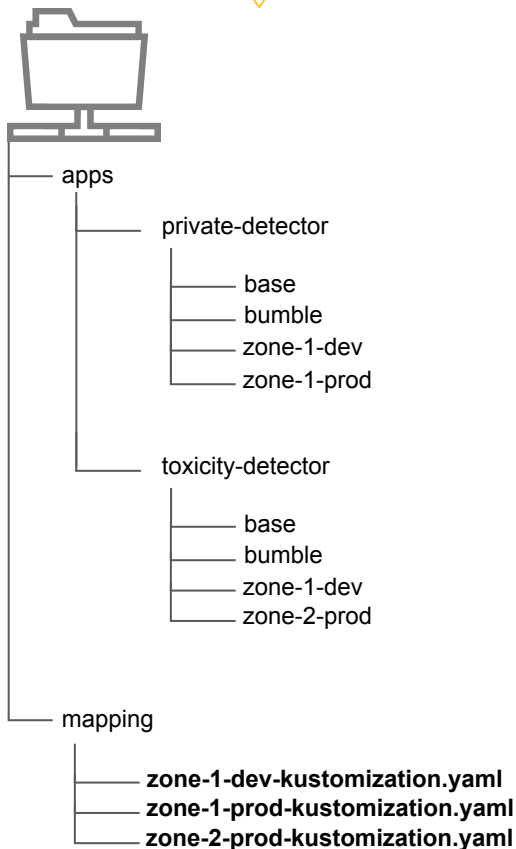
zone-1-dev complete manifest

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: private-detector
spec:
  replicas: 1
  template:
    metadata:
      labels:
        app: pd
    spec:
      containers:
        - name: pd
          image: pd:1.13.2
          ports:
            - containerPort: 80
          resources:
            limits:
              nvidia.com/mig-1g.5gb: 1
```

zone-1-prod complete manifest

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: private-detector
spec:
  replicas: 10
  template:
    metadata:
      labels:
        app: pd
    spec:
      containers:
        - name: pd
          image: pd:1.13.2
          ports:
            - containerPort: 80
          resources:
            limits:
              nvidia.com/gpu: 1
```

GitOps: Transparency



- **Transparency:** Each cluster has a list of the resources intended to run there.

mapping/zone-1-dev-kustomization.yaml

```
kind: Kustomization
resources:
- apps/private-detector/zone-1-dev
- apps/toxicity-detector/zone-1-dev
```

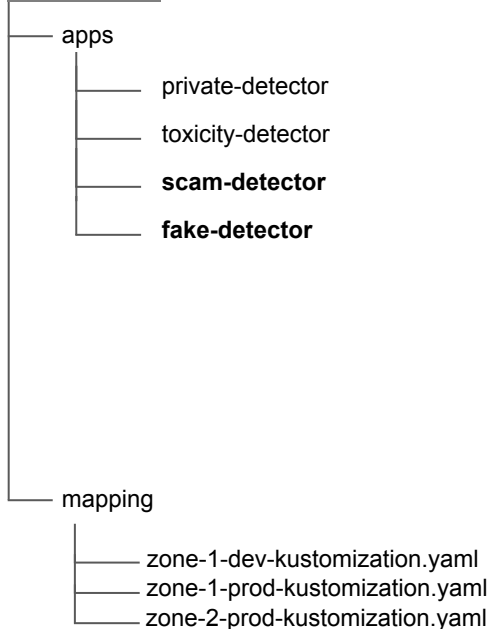
mapping/zone-1-prod-kustomization.yaml

```
kind: Kustomization
resources:
- apps/private-detector/zone-1-prod
# - apps/toxicity-detector/zone-1-prod
```

mapping/zone-2-prod-kustomization.yaml

```
kind: Kustomization
resources:
# - apps/private-detector/zone-2-prod
- apps/toxicity-detector/zone-2-prod
```


GitOps: Extensibility

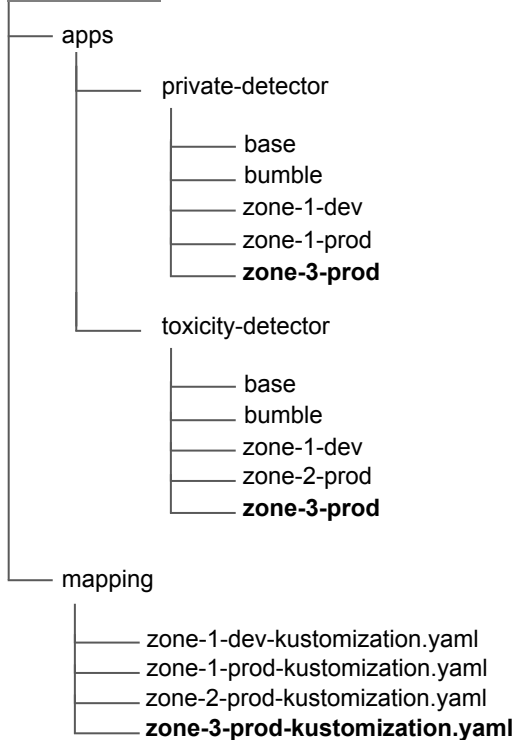


- **Transparency:** The process of adding new resources is simple and clear

mapping/zone-1-dev-kustomization.yaml

```
kind: Kustomization
resources:
- apps/private-detector/zone-1-dev
- apps/toxicity-detector/zone-1-dev
- apps/scam-detector/zone-1-dev
- apps/fake-detector/zone-1-dev
```

GitOps: Scalability

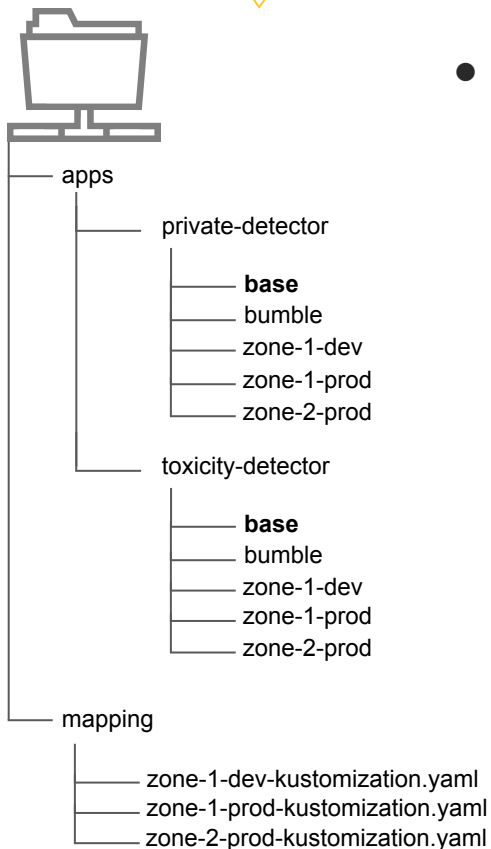


- **Scalability:** New environment could be added easily

mapping/zone-3-prod-kustomization.yaml

```
kind: Kustomization
resources:
- apps/private-detector/zone-3-prod
- apps/toxicity-detector/zone-3-prod
```

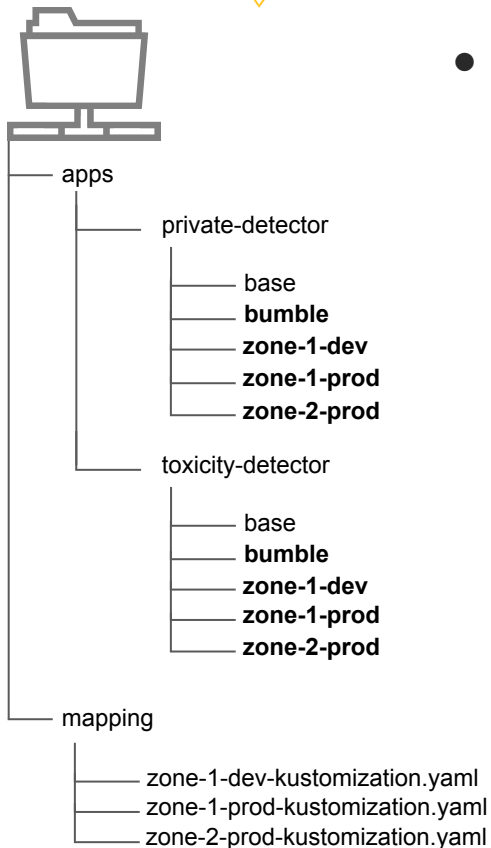
GitOps: Convenience



- **Convenience:** Resource upgrades require a single commit to be rolled out to all overlay clusters; overlay-specific parameters are declaratively defined.

GitOps: Collaboration

- **Collaboration:** Both cluster-specific and cluster-agnostic changes to a common base are described in a declarative way.



Thank you!